

Heat Transfer in a Metal Plate

1. Summary:

In this exercise you are asked to implement a GPU-accelerated application which simulates the heat propagation on a conductive metal plate. The progress of the simulation will be displayed through a 2D thermal imaging animation. The exercise aims to familiarize you with:

- Converting a studied physical phenomenon into a numerical simulation
- Using a GPU to speed up a parallel computing application
- Implementing a load balancer to split work between the CPU and the GPU for additional performance
- Using interoperability between GPU computing and graphics APIs to display the results of the numerical simulation

2. Physical Context:

Consider an environment which contains a 2D heat-conductive metal plate. The initial temperature of the plate and the surrounding environment (air) is T_R . Then we introduce a very small but constant heat source which provides a constant temperature T_S . The heat source is placed on the plate and starts warming it up. Because of the plate's conduction, the heat will progressively spread on the surface, as caused by the induced motion of atoms: the atoms closest to the heat source receive thermal energy and begin to vibrate, triggering collisions with other atoms; energy passed from these collisions makes other atoms vibrate and so on.

The heating of the plate is a process that begins with the introduction of the heat source and stops when the system reaches a thermal equilibrium (a state where there are no more variations in temperature in any point of the metal plate). In more technical terms, the process is called *transient heat conduction* and its equilibrium state at the end is called *steady-state conduction* (read more about both [here](#)). Throughout the time, the process is governed by the *heat equation* (read about it [here](#)).

3. Simulation details:

To simulate the heat transfer process described above, we will consider the metal plate to be a matrix of size $M \times N$. Given a moment of time t we consider $P[x, y](t)$ to represent the (instantaneous) heat of point (x, y) in that moment (with the constraints $0 \leq x < M$ and $0 \leq y < N$).

To simulate the initial condition, we consider the heat source to be placed at some given point (x_0, y_0) , meaning that $P[x_0, y_0](t) = T_s$ (constant over time). To model the surrounding environment, we considered the matrix to be “bordered” by one additional row or column on each side, holding only the value T_R in each point (also constant over time).

Given the cause of heat propagation as explained in the previous chapter, at any given moment in time t we will consider that the heat of any point is affected only by the heat values of its 8 direct neighboring points. For the points of the edge of the matrix (where x is 0 or $M - 1$ and where y is 0 or $N - 1$) some neighbor points will be on the “borders” (always having the constant value T_R).

Mean Filter Example

- Given a 3×3 mask, take the average of all nine values within the mask

```
>> x=uint8(10*magic(5))  
x =  
170 240 10 80 150  
230 50 70 140 160  
40 60 130 200 220  
100 120 190 210 30  
110 180 250 20 90
```

a	b	c
d	e	f
g	h	i

$$\rightarrow \frac{1}{9}(a + b + c + d + e + f + g + h + i)$$

The result of filtering x with 3×3 averaging filter

```
>> mean2(x(1:3,1:3))  
ans =  
111.1111  
110.0000  
131.1111
```

111.1111	108.8889	128.8889
110.0000	130.0000	150.0000
131.1111	151.1111	148.8889

6

Considering the state of the matrix P at some moment t , we now need to determine its state at the next moment, $t + 1$. For this we will use a simple 3×3 *mean filter* (an image processing example is given [here](#)). The filter will be applied independently to each point of matrix $P(t)$ in order to generate matrix $P(t + 1)$. For the application,

consider a point $P[x, y](t)$ and all of its 8 direct neighbors. Application of the mean filter (which simply calculates the average of all 9 points) is given by the following (intuitive) formula:

$$P[x, y](t + 1) = \frac{1}{9} \sum_{j=-1}^1 \sum_{i=-1}^1 P[x + j, y + i](t)$$

The earlier image provides a numerical example (you can read more about it [here](#)). It is now clear to see that calculating $P(t + 1)$ from $P(t)$ can be done in parallel with ease (namely, any point $P[x_1, y_1](t)$ can be calculated in parallel to any other point $P[x_2, y_2](t)$) – in fact, this is an example of an *embarrassingly parallel* application.

The heat transfer simulation effectively starts with $P(t = 0)$ and proceeds to calculate $P(t = 1)$, then $P(t = 2)$ and so on, until a *state of convergence* is reached. This state corresponds to the thermal equilibrium state described in the previous chapter and occurs when there is no more significant change between $P(t + 1)$ and $P(t)$. Numerically, this means that the sums of all absolute differences between the new and current state of P fall below a (small) threshold:

$$\left(\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} |P[x, y](t + 1) - P[x, y](t)| \right) < \epsilon$$

The initial state corresponds to the metal plate sitting in ambient air without the heat source being applied; therefore, in this state the matrix $P(t = 0)$ contains the value T_R in all points.

4. Deliverables and Scoring:

You are asked to implement the following tasks:

- 1) Take as input M, N (the dimensions of the temperature matrix for the metal plate), x_0, y_0 (the point where the heat source is placed), T_S and T_R . Run the simulation by calculating $P(t)$ (namely, $P(1)$, then $P(2)$ and so on) as explained earlier *in parallel* using a GPU computing device of your choice and your preferred GPU computing API (OpenCL or CUDA). Provide a notification when convergence is reached. **(45 points)**

- 2) To visualize the simulation, you will need to display $P(t)$ as a 2D thermal field image (an example is [here](#)). To convert a temperature value to RGB for the display you may use the algorithm found [here](#), or any other alternative of your choice. To generate the image, you must use the interoperability with OpenGL or DirectX available in the OpenCL or CUDA APIs. **(30 points)**
- 3) Implement a load balancing mechanism, where a percentage $F\%$ of values is calculated on the GPU and the remaining percentage of $(100 - F)\%$ is calculated on the CPU in parallel (using `pthread`s or any other API of your choice). The value of F is given as input ($0 \leq F \leq 100$) and may be changed dynamically by the user throughout the simulation. To see the effects of when balancing the computation load you will display the average time needed to calculate a single simulation step (that is, full calculation of $P(t)$ for a given moment t). The time will be displayed in the console or together with the RGB image, depending on your choice. **(15 points)**
- 4) Allow the user to change the problem configuration dynamically (i.e. the users may change the position of the heat source by providing a different position (x_0, y_0) , they may change the heat source temperature T_s or the air temperature T_R , however the matrix dimensions M and N will remain constant throughout the simulation). The simulation should not be restarted from the beginning when changing the configuration. **(10 points)**